

An aerial night view of the Chicago skyline, showing numerous illuminated buildings and streets, creating a bokeh effect in the background.

OPERATIONALIZING R WITH ORACLE AND POSTGRES

EXAMPLES OF DATA SCIENCE PROJECTS
AT THE CITY OF CHICAGO

PRESENTED TO CHICAGO R USER GROUP,
AUGUST 30, 2018

Example Projects

Food Inspection Forecasting



Clear Water




Lead Safe




Vector Borne Diseases



Project Pattern



Project Pattern



RESOURCES

RSTUDIO

- RStudio has released a wiki for Database connection info:

<https://db.rstudio.com/odbc/>

The screenshot shows the RStudio Database Wiki page for ODBC. The page has a navigation bar at the top with links: Overview, Getting Started, Best Practices, Databases, and Blog. A search icon is on the right. On the left, there is a sidebar with a search bar and a list of categories: Getting Started, Packages, and Database Queries. The main content area is titled 'odbc' and contains a paragraph explaining the goal of the 'odbc' package: to provide a DBI-compliant interface to Open Database Connectivity (ODBC) drivers. It lists supported databases: SQL Server, Oracle, MySQL, PostgreSQL, SQLite, and others. It also mentions that the implementation builds on the 'nanodbc' C++ library. To the right of the text is a diagram titled 'Commercial Databases' showing the flow from 'Data' (Tables & Views) through a 'Database Driver' (SQL Engine) to 'Access & Wrangle' (odbc package, DBI package, dplyr package) and finally to 'R Studio'. Below the diagram is a 'Usage' section with a paragraph stating that all examples assume a connection 'con' has been created and that users should visit the 'Databases' page for more information.

Databases using R from RStudio

Overview Getting Started Best Practices Databases Blog

Getting Started

Connect to a Database

Database Queries

Packages

dplyr

DBI

odbc

pool

dbplot

odbc

The goal of the `odbc` package is to provide a DBI-compliant interface to [Open Database Connectivity](#) (ODBC) drivers. This allows for an efficient, easy to setup connection to any database with ODBC drivers available, including [SQL Server](#), [Oracle](#), [MySQL](#), [PostgreSQL](#), [SQLite](#) and others. The implementation builds on the [nanodbc](#) C++ library.

Usage

All of the following examples assume you have already created a connection `con`. To find out how to connect to your specific database type, please visit the [Databases](#) page.

Commercial Databases

Data	Access & Wrangle
Tables & Views	<code>odbc</code> package
SQL Engine	<code>DBI</code> package
Database	<code>dplyr</code> package

R Studio

RODBC / Brian Ripley

- Ripley's RODBC vignette is still my first reference, and goes way beyond ODBC:

<https://cran.r-project.org/web/packages/RODBC/vignettes/RODBC.pdf>

ODBC Connectivity

by Brian Ripley
Department of Statistics, University of Oxford
ripley@stats.ox.ac.uk

May 5, 2017

Package **RODBC** implements ODBC database connectivity. It was originally written by Michael Lapsley (St George's Medical School, University of London) in the early days of R (1999), but after he disappeared in 2002, it was rescued and since much extended by Brian Ripley. Version 1.0-1 was released in January 2003, and **RODBC** is nowadays a mature and much-used platform for interfacing R to database systems.

Thanks to Marc Schwartz for contributing some of the experiences here. See also the archives of the **R-sig-db** mailing list.

BASICS DATABASES WITHIN R

DB PACKAGES

- **RODBC** – This is the original R database connector, written by Brian Ripley and released in 2003
- **ROracle** – Released by Oracle and rarely maintained, but it does work
- **RPostgreSQL** – New, well documented, great dev team
- **mongolite** – This is our default MongoDB package
- **DBI** – Not used directly, but it's the workhorse behind 2 & 3

BASIC CONCEPTS

- Create a connection to the driver or named ODBC connection
- Execute commands against the connection
 - Commands can be meta (e.g. “ListTables”)
 - Commands can be SQL (e.g. “sendQuery”)
 - Commands can use a “cursor” that allows you to “page” through your response


EXAMPLE - CONNECTING

- Create a “channel” or “connection” object
- The connection is typically named “ch” or “con”
- Meta commands like `odbcGetInfo` provide information about the connection

```
## LOAD LIBRARY
library("RODBC")
library("data.table")

## CREATE CONNECTION OBJECT / CONNECT TO DATABASE
con <- odbcConnect(dsn = "BUILDING_INSPECTIONS",
                  uid = "data_science_team",
                  pwd = "bfb60492126e&pf")


odbcGetInfo(con)
```



EXAMPLE – TYPICAL QUERY

- There are meta commands like “get table” that are basically “select star” commands
- You can put nearly any query into functions that submit generic queries
- The package attempts to format the return, sometimes there are issues

```
dat <- sqlQuery(channel = con,  
               query = paste("select * FROM Example_Inspections",  
                             "WHERE status = 'COMPLETE'"),  
               stringsAsFactors = FALSE)
```



NOTE: stringsAsFactors
= FALSE can greatly
improve
performance!

EXAMPLE – EXPLORING THE DB

- There are many ways to find out what's in the database
 - Meta commands like “sqlTables” will tell you tables
 - System specific commands will also provide information on the tables
- Everything uses the “connection” object to access the database via the driver

ORACLE SYNTAX

```
## GET ALL TABLE NAMES
all_tables <- sqlTables(con, tableType = "TABLE")

## GET COLUMNS ATTRIBUTES FROM ALL TABLES
all_columns <- data.table(sqlQuery(
  channel = con,
  query = paste("SELECT TABLE_CATALOG, TABLE_SCHEMA, TABLE_NAME,",
                "ORDINAL_POSITION, COLUMN_NAME, COLUMN_DEFAULT, IS_NULLABLE,",
                "DATA_TYPE",
                "FROM INFORMATION_SCHEMA.COLUMNS",
                "ORDER BY TABLE_NAME, COLUMN_NAME"),
  stringsAsFactors = FALSE))
```

POSTGRES SYNTAX

```
schemas <- as.data.table(dbGetQuery(con,
                                     "SELECT * FROM information_schema.tables"))
schemas[grepl("wic", table_name)]

sort(schemas[grepl("aux", table_schema), table_name])
```

PACKAGE METHODS

```
library("ROracle")
```

```
ROracle::
```

◆ dbHasCompleted	{ROracle}
◆ dbListConnections	{ROracle}
◆ dbListFields	{ROracle}
◆ dbListResults	{ROracle}
◆ dbListTables	{ROracle}
◆ dbReadTable	{ROracle}
◆ dbRemoveTable	{ROracle}
◆ dbRollback	{ROracle}

dbListFields

These functions mimic their R counterparts except that they generate code that gets remotely executed in a database engine: get, assign, exists, remove, objects, and names.

Press F1 for additional help

- Many packages (RODBC, ROracle, RPostgreSQL) have built in convenience functions that *mimic* SQL functions
- You can use these to read / write tables
- They also have something that “just executes” a query, e.g. RPostgreSQL

CODE ORGANIZATION


INLINE QUERIES

- Simple strategy, works well if you only do SELECT *
- Gets messy with bigger queries

```
dat <- sqlQuery(channel = con,  
               query = paste("select * FROM Example_Insp",  
                             "WHERE status = 'COMPLETE'"),  
               stringsAsFactors = FALSE)
```

TWO STRATEGIES FOR KEEPING THE CODE PUBLIC

Option 1: Totally Public



Clear Water



- Uses RSocrata to leverage the Socrata data portal
- Everything stays public
- Nothing to worry about with Postgres or Oracle

Option 2: YAML Files

- Put all the code online
- Keep connection info in plain text in one folder (like ./params)
- Add folder to .gitignore
- Bonus: Use YAML files
 - More readable
 - Named fields
 - Only one dependency

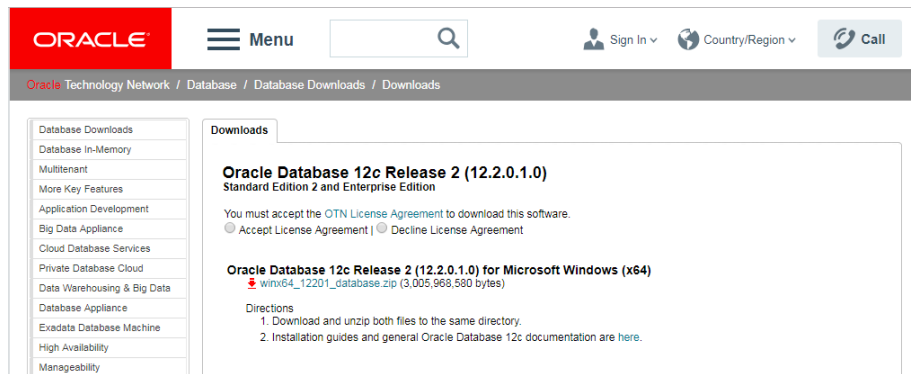
```
## LOAD PARAMETER DATA
params <- yaml::yaml.load_file(param_file_loc)
dbinfo <- yaml::yaml.load_file(dbinfo_file_loc)

## CREATE DATABASE CONNECTION TO ORACLE DATABASE
drv <- DBI::dbDriver("Oracle")
connect_string <- paste0(
  "(DESCRIPTION=",
  "(ADDRESS=(PROTOCOL=tcp)(HOST=", dbinfo$host, ")(PORT=1521))",
  "(CONNECT_DATA=(SERVICE_NAME=", dbinfo$service, ")))")
ch <- ROracle::dbConnect(drv,
  username = dbinfo$username,
  password = dbinfo$password,
  dbname = connect_string)
```

WINDOWS CLIENT ENVIRONMENT

INSTALL DRIVERS

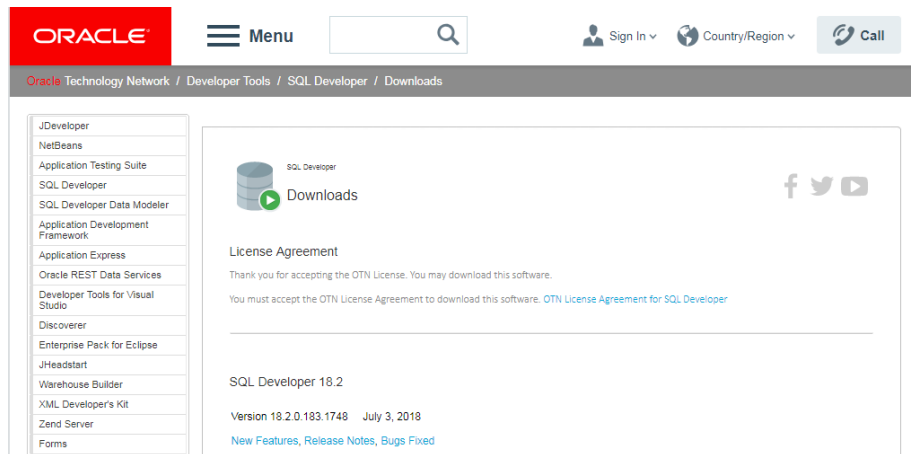
- Download drivers from Oracle
- Main flavors
 - 11g
 - 12c
- Many driver options, I used trial and error to select correct driver



INSTALL SQL DEVELOPER

(Optional)

- It's very helpful to have SQL Developer for checking connections /testing queries

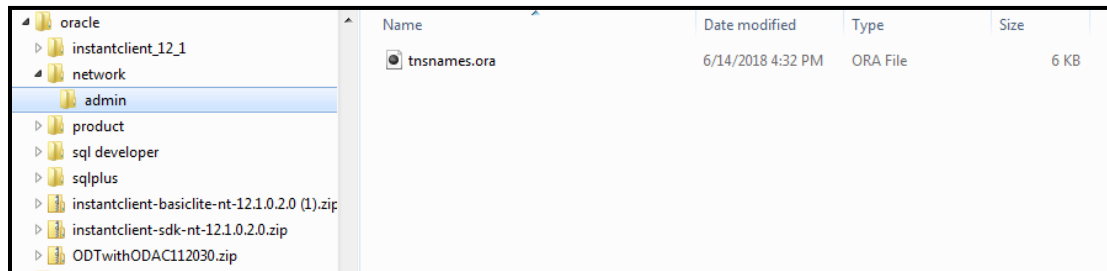


SET UP TNS NAMES

- The sane location for your Oracle install is c:\oracle
- Add connection info to tnsnames.ora

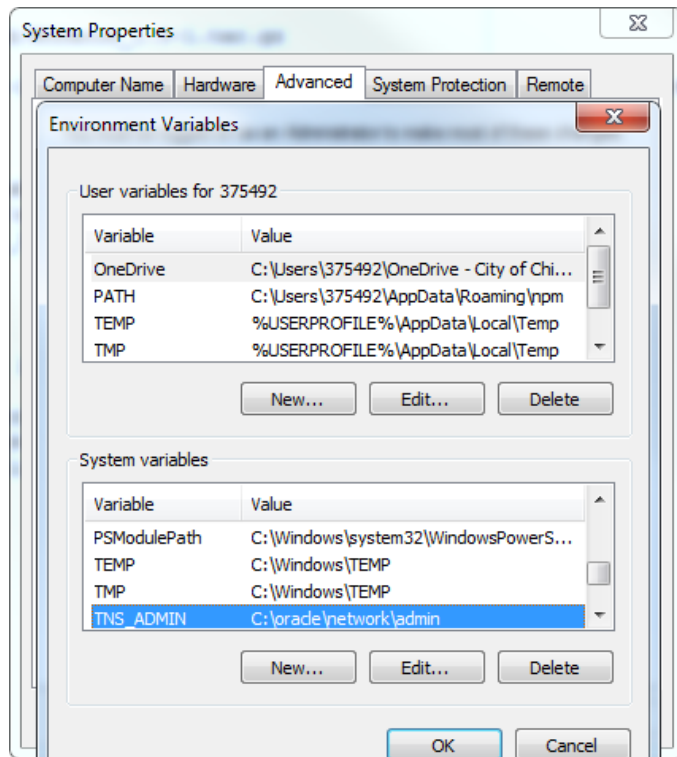
```
## CONNECTION NOTES FOR THIS SERVER
## THIS IS AN EXAMPLE TNS ENTRY
## SEE EMAIL ON 4/9/15
## Username: SomeoneSpecial
## Password: DoIT4Life

GIS_SERVER =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP) (HOST = PX02-WORK) (PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = GISDEV)
    )
  )
```



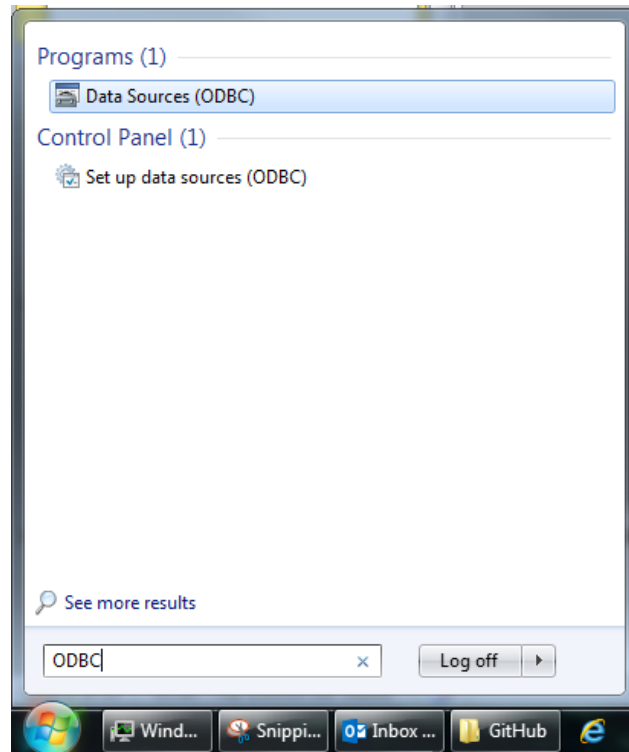
ADD TNS_ADMIN

- Set an **environmental variable** for the TNS File
- Needed for
 - SQL Developer
 - ODBC management
 - ROracle package
- At this point you can add TNS connections by name in SQL Developer
 - Look for a green plus sign
 - Choose TNS for connection type
 - Named connections appear in dropdown



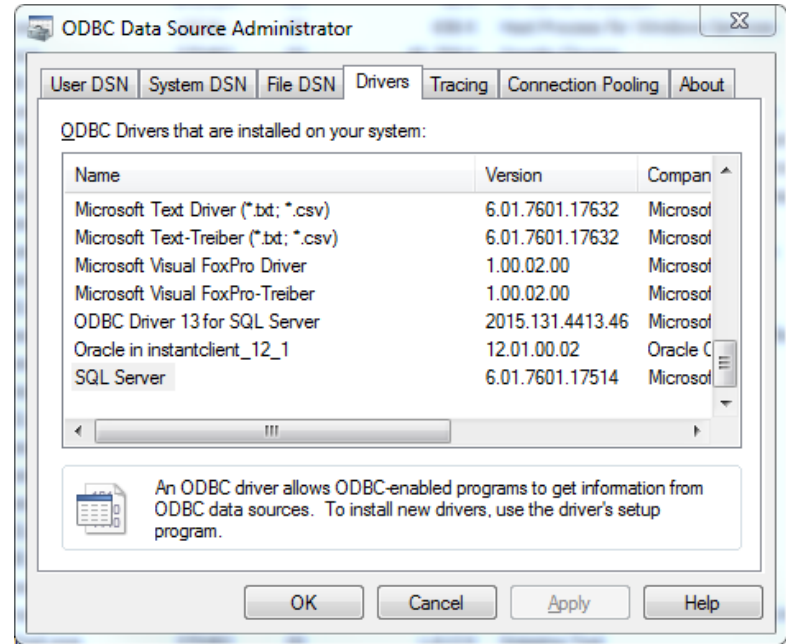
ODBC CONNECTION

- Windows has a “data sources” manager
- Search for ODBC



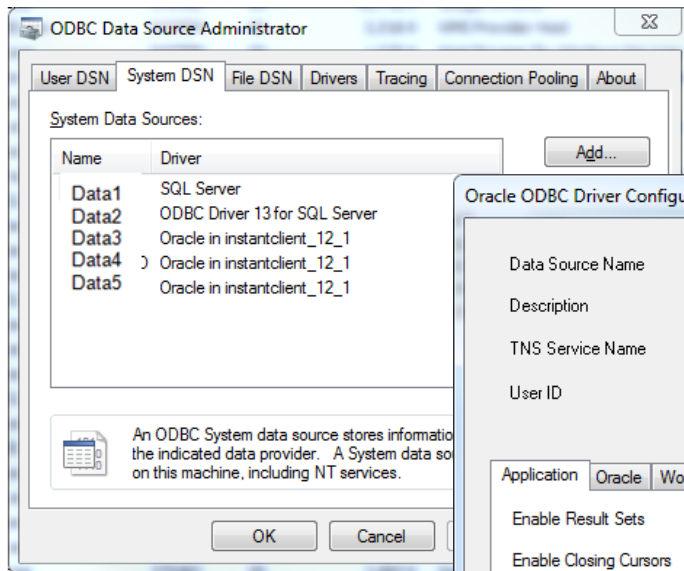
ODBC CONNECTION

- Default drivers are listed
- Your newly installed driver is listed here

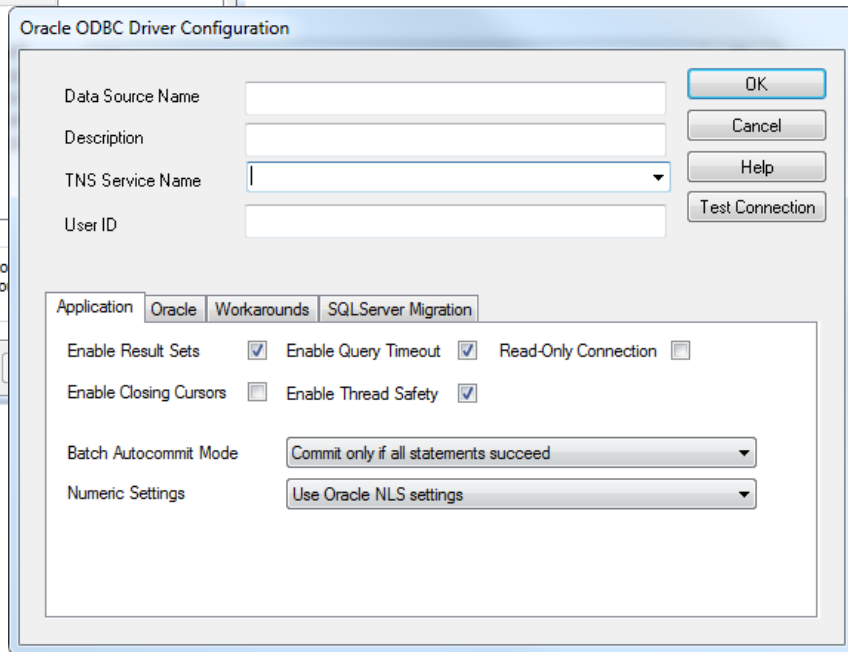


ADD SYSTEM DSN

- This is a list of your data source names
- Click “Add...” and follow the prompts



- This menu appears, and (for Oracle) your TNS names are available here



SPECIFY TNS WITHIN R

- With the ROracle package, you can pass in the TNS entry
- Pros:
 - Simpler?
 - Code is more portable
 - SQL Developer is configured by default
- Cons:
 - ROracle is hard to install
 - Simpler to refer to a named entry
 - SQL Developer isn't configured

```
## LOAD PARAMETER DATA
params <- yaml::yaml.load_file(param_file_loc)
dbinfo <- yaml::yaml.load_file(dbinfo_file_loc)

## CREATE DATABASE CONNECTION TO ORACLE DATABASE
drv <- DBI::dbDriver("Oracle")
connect_string <- paste0(
  "(DESCRIPTION=",
  "(ADDRESS=(PROTOCOL=tcp)(HOST=", dbinfo$host, ")(PORT=1521))",
  "(CONNECT_DATA=(SERVICE_NAME=", dbinfo$service, ")))")
ch <- ROracle::dbConnect(drv,
  username = dbinfo$username,
  password = dbinfo$password,
  dbname = connect_string)
```

ROracle

ROracle IN WINDOWS

- Oracle SDK is required
- SDK Note: you just unzip the files to the oracle folder and mingle them with other files that are already there
- More environmental variables are required (see picture on right)
- Must be installed at the command line, not within R*

```
## WINDOWS INSTALLATION
## Download SDK unzip to oracle folder
## Create the following environmental variables:
##     OCI_INC=C:\oracle\instantclient_12_1\sdk\include
##     OCI_LIB32=C:\oracle\instantclient_12_1
## Then run:
R CMD INSTALL ROracle 1.3-1.tar.gz
```

* I'm pretty sure of this, but not positive

ROracle IN LINUX

- Install Oracle drivers

```
cd /app/Oracle/  
  
unzip instantclient-basiclite-linux-11.2.0.4.0.zip  
unzip instantclient-odbc-linux32-11.2.0.2.0.zip  
unzip instantclient-sqlplus-linux32-11.2.0.2.0.zip  
  
cd instantclient_11_2  
ln -s libclntsh.so.11.1 libclntsh.so  
ln -s libocci.so.11.1 libocci.so
```


ROracle IN LINUX

- Create a script to set the appropriate paths
- I added my script to my ~/.bashrc

```
export ORACLE_HOME=/app/Oracle/instantclient_11_2
export LD_LIBRARY_PATH=$ORACLE_HOME
export TNS_ADMIN=/app/Oracle
export PATH=$ORACLE_HOME:$PATH:$LD_LIBRARY_PATH
export SERVICE_NAME=GIS_DATABASE

# export LD_LIBRARY_PATH=/app/Oracle/instantclient_11_2:$LD_LIBRARY_PATH
# export PATH=/app/Oracle/instantclient_11_2:$PATH

echo $(printenv | grep ORACLE)
echo $(printenv | grep LD_LIBRARY_PATH)
echo $(printenv | grep TNS_ADMIN)
echo $(printenv | grep PATH)
```

ROracle IN LINUX

- Install ROracle
- Note: I'm setting the environmental variables with set_oracle_env.sh

```
## LINUX INSTALLATION
cd ~
wget https://cran.r-project.org/src/contrib/ROracle_1.3-1.tar.gz
source /app/Oracle/set_oracle_env.sh
R CMD INSTALL \
    --configure-args='--with-oci-lib=/usr/lib/oracle/12.1/client64/lib \
    --with-oci-inc=/usr/include/oracle/12.1/client64' ROracle_1.3-1.tar.gz
```

TIME IN ORACLE

TIME ZONES

- The West Nile Prediction repository has full examples (url below)
- Note: you can get and set env variables from within R
- Note: You need to set two variables in order to upload date / times

<https://github.com/Chicago/west-nile-virus-predictions/>

```
#####  
## LOG INTO ORACLE (GET USER INFO FROM TNS FILE)  
#####  
  
dbinfo_dev <- readLines("untracked/zdt_credentials_dev.txt")  
dbinfo_prod <- readLines("untracked/zdt_credentials_prod.txt")  
# dbinfo_user <- readLines("untracked/zdt_credentials.txt")  
  
drv <- dbDriver("Oracle")  
# system("echo $TNS_ADMIN")  
# Sys.getenv("TNS_ADMIN")  
# Sys.setenv(TNS_ADMIN = "/app/oracle")  
  
Sys.setenv(TZ = "GMT")  
Sys.setenv(ORA_SDTZ = "GMT")  
  
connect_string_dev <- paste0(  
  "(DESCRIPTION=",  
  "(ADDRESS=(PROTOCOL=tcp)(HOST=", dbinfo_dev[1], ")(PORT=1521))",  
  "(CONNECT_DATA=(SERVICE_NAME=", dbinfo_dev[2], ")))"  
)  
ch_dev <- ROracle::dbConnect(drv, username = dbinfo_dev[3],  
  password = dbinfo_dev[4],  
  dbname = connect_string_dev)  
  
connect_string_prod <- paste0(  
  "(DESCRIPTION=",  
  "(ADDRESS=(PROTOCOL=tcp)(HOST=", dbinfo_prod[1], ")(PORT=1521))",  
  "(CONNECT_DATA=(SERVICE_NAME=", dbinfo_prod[2], ")))"  
)  
ch_prod <- ROracle::dbConnect(drv, username = dbinfo_prod[3],  
  password = dbinfo_prod[4],  
  dbname = connect_string_prod)
```

RJDBC

RJDBC – THE PROMISE OF JAVA

- RJDBC is a lot easier
- I've had issues in Linux (user error?)
- With RJDBC the setup should be simpler and platform independent
- This simple example (for Windows) should work in Linux

```
library(RJDBC)
username = paste(readLines("username.txt"))
password = paste(readLines("password.txt"))
drv <- JDBC("oracle.jdbc.driver.OracleDriver",
            "C:/Program Files/Oracle/ojdbc7.jar")
violation_conn <- dbConnect(drv,
                             "jdbc:oracle:thin:@192.168.0.166:1521/DBNAME",
                             username,
                             password)
q <- paste(readLines("queries/inspection_data.sql"), collapse="")
insp <- dbGetQuery(con, statement = q)
```

To be continued...

THANK YOU

Contact Info:

Gene Leynes

Data Scientist @ City of Chicago



Email: datascience@cityofchicago.org

Websites:

<https://github.com/Chicago/west-nile-virus-predictions>